

# MVA – Probabilistic Graphical Models

## Homework 2

Wilson JALLET\*

December 12, 2019

### 1 K-means and the EM algorithm

**Question 1.** We consider a mixture model of  $K$  components for a dataset  $(X_i)$  where  $X_i \in \mathbb{R}^d$ . We denote  $Z_i \in \{1, \dots, K\}$  the latent hidden label. Each component occurs with probability  $p_k = \mathbb{P}(Z_i = k)$  and is distributed as

$$X_i \sim \mathcal{N}(\mu_k, D_k)$$

i.e.  $p(x|k) = \frac{1}{((2\pi)^d |D_k|)^{1/2}} \exp(-\frac{1}{2}(x - \mu_k)^T D_k^{-1} (x - \mu_k))$ .

The data log-likelihood under parameters  $\Theta = ((p_1, \mu_1, D_1), \dots, (p_K, \mu_K, D_K))$  is

$$\mathcal{L}(X_1, \dots, X_n; \Theta) = \sum_{i=1}^n \log \left( \sum_{k=1}^K p_k p(X_i|k; \mu_k, D_k) \right)$$

We seek to compute the MLE

$$\hat{\Theta} \in \operatorname{argmax}_{\Theta} \mathcal{L}(X_1, \dots, X_n; \Theta)$$

This optimization problem is intractable when using straightforward methods.

The EM algorithm goes as follows:

- Expectation Compute the posterior probability of the latent variables  $Z_i$ :

$$q_{k,i}^{(t)} = p(Z_i = k | X_i; \Theta^{(t)}) = \frac{p_k^{(t)} p(X_i | Z_i = k; \Theta^{(t)})}{\sum_{\ell=1}^K p_{\ell}^{(t)} p(X_i | Z_i = \ell; \Theta^{(t)})} \quad (1)$$

and denote  $w_k^{(t)} = \sum_{i=1}^n q_{k,i}^{(t)}$  – we then have  $\sum_{k=1}^K w_k^{(t)} = n$ .

- Maximization Update the parameters  $\Theta^{(t)}$  by maximizing the lower bound objective:

$$\max_{\Theta} \mathcal{J}(q^{(t)}, \Theta) = \sum_{i=1}^n \left( \sum_{k=1}^K q_{k,i}^{(t)} \left( \log p_k - \frac{d}{2} \log(2\pi) - \frac{1}{2} \log |D_k| - \frac{1}{2} (X_i - \mu_k)^T D_k^{-1} (X_i - \mu_k) \right) \right) \quad (2)$$

---

\*[wilson.jallet@polytechnique.org](mailto:wilson.jallet@polytechnique.org)

subject to  $\sum_{k=1}^K p_k = 1$  (associated to a multiplier  $\nu$ ). The KKT conditions give the null gradient condition:

$$\frac{1}{p_k} w_k^{(t)} - \nu = 0 \quad (3a)$$

$$-\sum_{i=1}^n q_{k,i}^{(t)} D_k^{-1} (\mu_k - X_i) = 0 \quad (3b)$$

$$-\frac{1}{2} \sum_{i=1}^n q_{k,i}^{(t)} (D_k^{-1} - D_k^{-2} \text{diag}(X_i - \mu_k)^2) = 0 \quad (3c)$$

Which leads to the updates:

$$p_k = \frac{1}{n} w_k^{(t)} \quad (4a)$$


$$\mu_k = \frac{1}{w_k^{(t)}} \sum_{i=1}^n q_{k,i}^{(t)} X_i \quad (4b)$$

$$D_k = \frac{1}{w_k^{(t)}} \sum_{i=1}^n q_{k,i}^{(t)} \text{diag}(X_i - \mu_k)^2 \quad (4c)$$

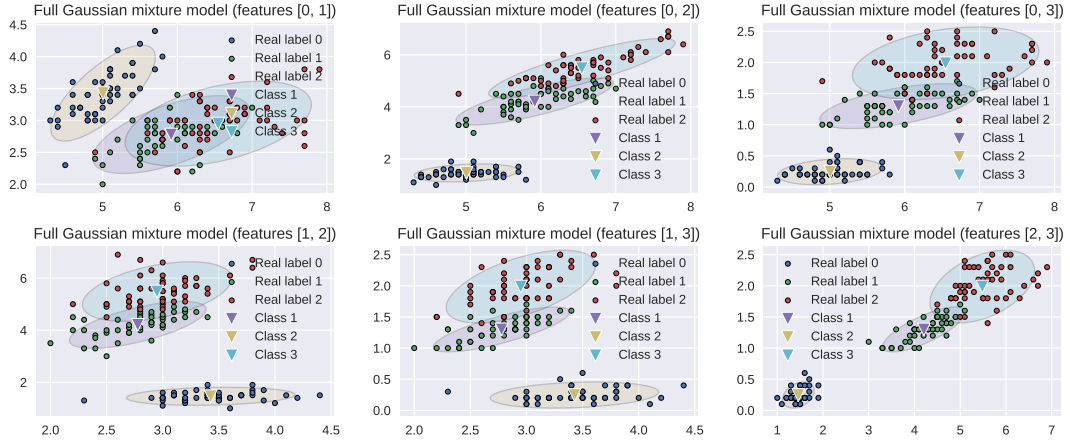
**Question 2.** The main advantage of this “reduced” covariance mixture model is that it is more sparse: it uses far fewer parameters ( $K(2d + 1)$ ) than the its full counterpart, which has  $K(1 + d + d(d + 1)/2)$  parameters. For datasets with relatively independent features (conditionally on the latent class), this can give performance very close to the full covariance while having a smaller, simpler model (meaning better AIC or BIC scores).

**Question 3.** Figure 1 compares the obtained latent class centroids and confidence ellipsoids (where applicable) for the diagonal and full covariance mixture models and K-means, on the Iris dataset, for a small number of classes  $K = 3$  (the actual number of classes in the data). Figures 2 and 3 represent the same for  $K = 2, 4$  classes.


**Question 4.** I created synthetic data made up of an ellipsis and annulus: Figure 4 shows how the mixture models and K-means compare.



(a) Diagonal Gaussian mixture model on the Iris dataset (our implementation).  $K = 3$  classes.

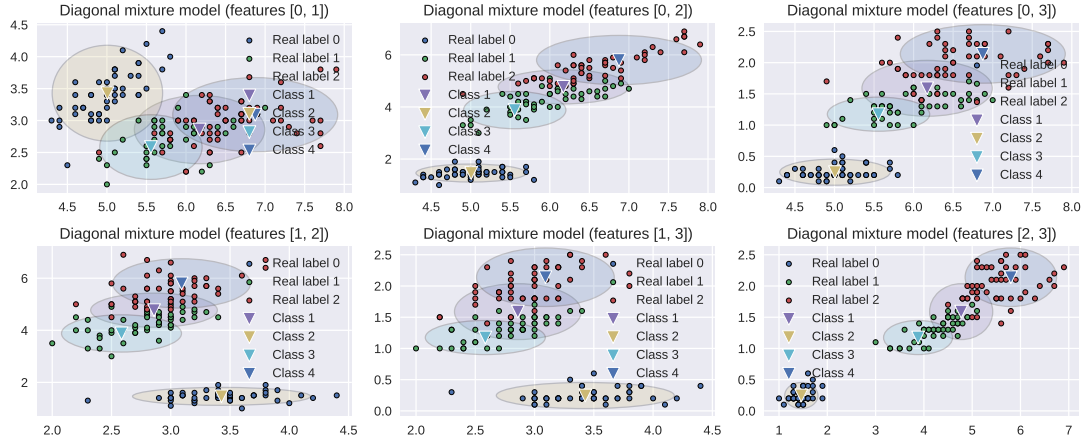


(b) Full Gaussian mixture on the Iris dataset using `scikit-learn`.  $K = 3$  classes.

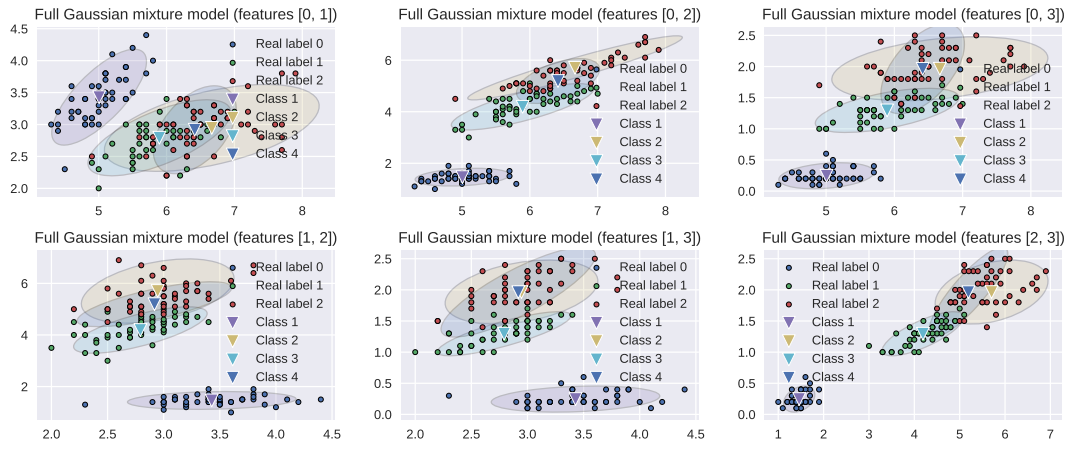


(c) Centroids of the K-means model.

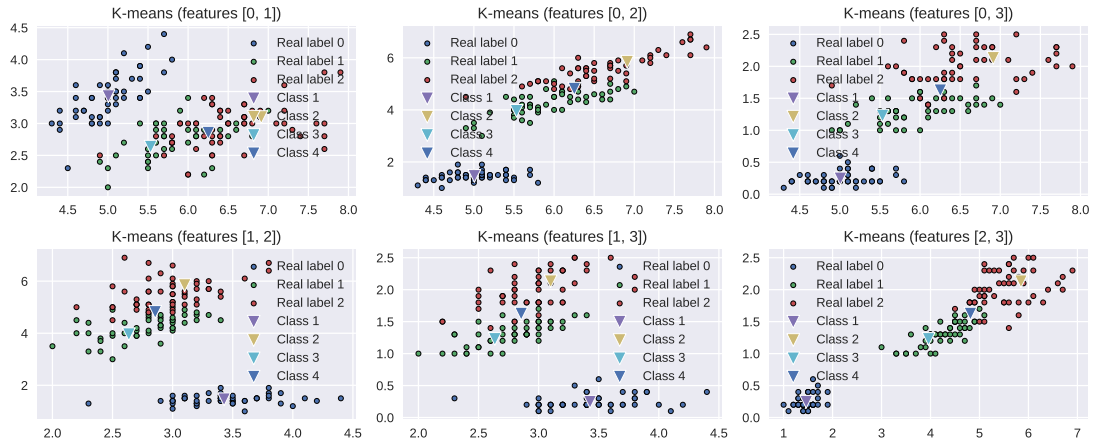
Figure 1: Comparison of the diagonal and full covariance mixture models and K-means for  $K = 3$  classes.



(a) Diagonal Gaussian mixture model on the Iris dataset (our implementation).

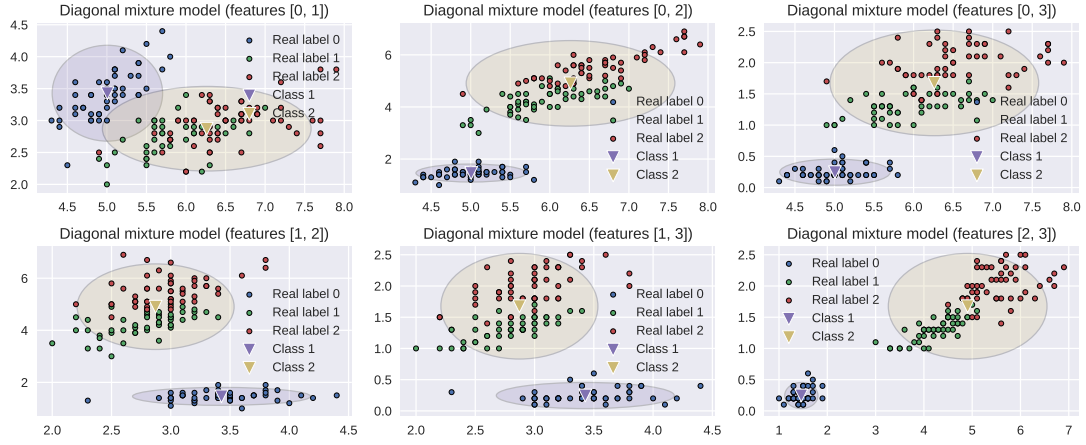


(b) Full Gaussian mixture on the Iris dataset using `scikit-learn`.

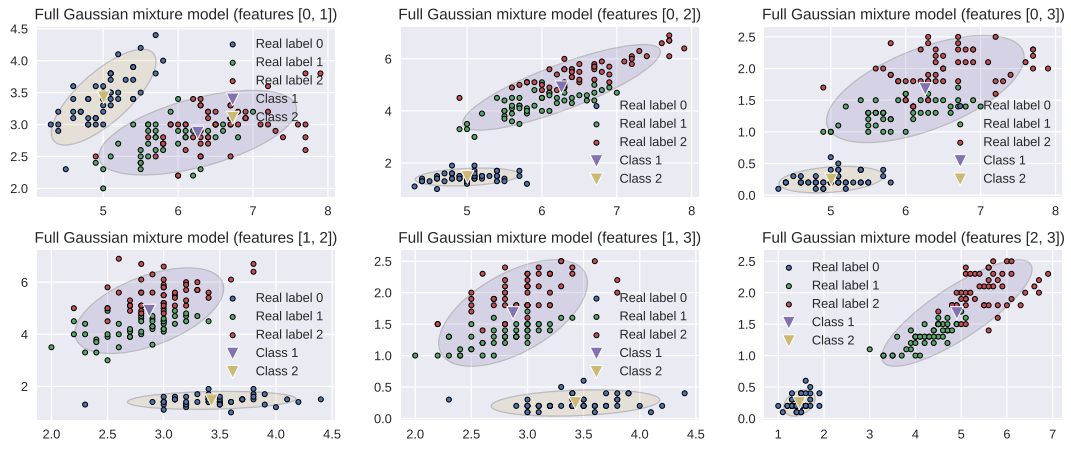


(c) K-means.

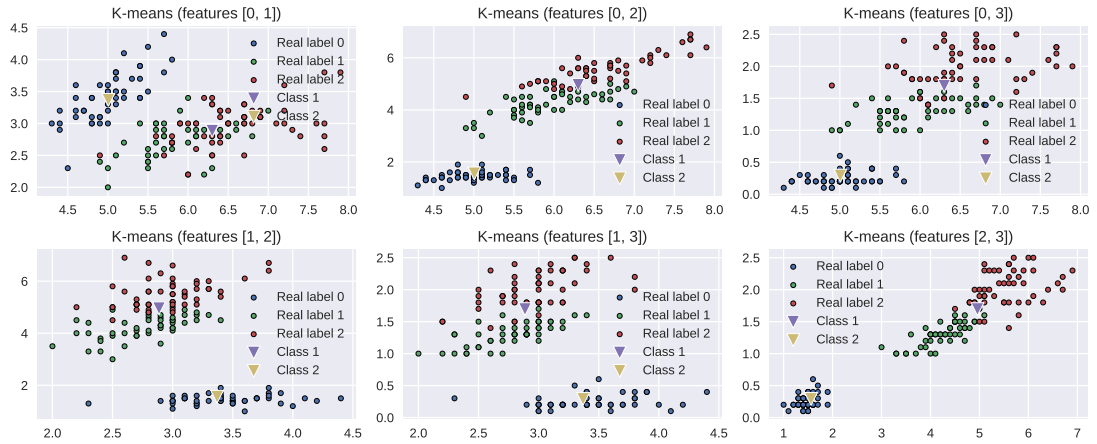
Figure 2: Comparison of the models for  $K = 4$  classes.



(a) Diagonal Gaussian mixture model on the Iris dataset (our implementation).




(b) Full Gaussian mixture on the Iris dataset using `scikit-learn`.




(c) Full Gaussian mixture on the Iris dataset using `scikit-learn`.


Figure 3: Comparison of the models for  $K = 2$  classes.



(a) Diagonal Gaussian mixture model.



(b) Full Gaussian mixture model.



(c) Result of the K-means model: notice how the class centroids are offset from the actual center of the class.

Figure 4: Comparison of the models on synthetic data where K-means performs worse than mixtures.

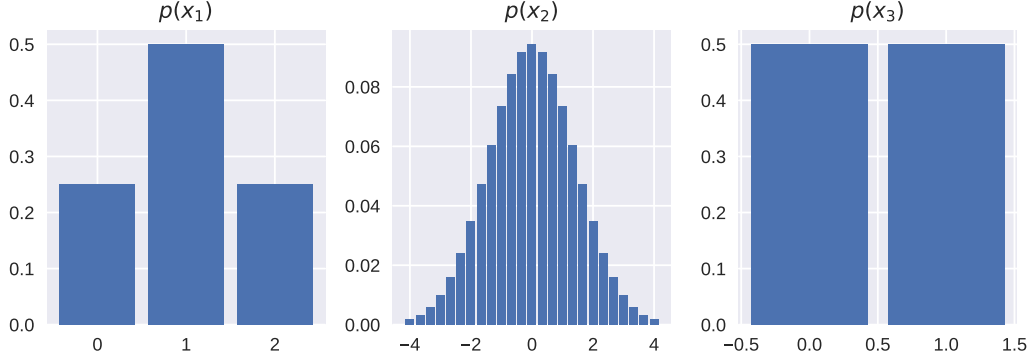


Figure 5: Sanity test: undirected graph with  $n = 3$  nodes, and a probability distribution  $p$  with factors  $\psi_1 = (1, 2, 1)$ ,  $\psi_2(x_2) = \exp(-x_2^2/4)$ ,  $\psi_3 = (1, 1)$  and  $\psi_{i,i+1} = 1$  (i.e. independence).

## 2 Graphs, algorithms and Ising

**Question 1.** We recall that for an undirected chain graph  $G$  probability distributions factor as

$$p(x) = \frac{1}{Z} \prod_{i=1}^n \psi_i(x_i) \prod_{i=1}^{n-1} \psi_{i,i+1}(x_i, x_{i+1}) \quad (5)$$

**The algorithm.** The marginal distribution of  $X_i$  can be rewritten as

$$p(x_i) = \frac{1}{Z} \mu_{i-1,i}(x_i) \psi_i(x_i) \mu_{i+1,i}(x_i)$$

where  $\mu_{i-1,i}, \mu_{i+1,i}$  are (forward, backward) messages from  $i-1$  to  $i$  and  $i+1$  to  $i$ . They are propagated as:

$$\mu_{i,i+1}(x_{i+1}) = \sum_{x_i} \psi_i(x_i) \psi_{i,i+1}(x_i, x_{i+1}) \mu_{i-1,i}(x_i) \quad (6a)$$

$$\mu_{i,i-1}(x_{i-1}) = \sum_{x_i} \psi_i(x_i) \psi_{i-1,i}(x_{i-1}, x_i) \mu_{i+1,i}(x_i) \quad (6b)$$

**Practical implementation.** If the state space  $\mathcal{X}$  of the variables  $X_1, \dots, X_n$  (for instance for binary variables) is discrete we can represent the input functions  $\psi_i$  and  $\psi_{i,i+1}$  as arrays. If not (continuous variables for instance), we can discretize a grid over  $\mathcal{X}$  and precompute an array of values for the factors. Denoting the arrays in bold letters, we forward-propagate by

$$\boldsymbol{\mu}_{i,i+1} = (\boldsymbol{\mu}_{i-1,i} \odot \boldsymbol{\psi}_i) \boldsymbol{\psi}_{i,i+1}$$

and back-propagate by

$$\boldsymbol{\mu}_{i,i-1} = (\boldsymbol{\mu}_{i,i+1} \odot \boldsymbol{\psi}_i) \boldsymbol{\psi}_{i-1,i}^T$$

which allows to compute the marginal distributions using vectorized operations.

Figure 5 shows an implementation for independent edges. The expected marginals are  $X_1 \sim \mathcal{M}(1; 1/4, 1/2, 1/4)$ ,  $X_2 \sim \mathcal{N}(0, 2)$  and  $X_3 \sim \mathcal{B}(1/2)$ .

**Question 2.** The vertex set  $V$  of the graph  $G = (V, E)$  induced by the grid has vertices of the form  $v = (j, k)$  where  $1 \leq j \leq w$  and  $1 \leq k \leq h$ . An easy junction tree to extract from  $G$  is given by collapsing the rows into supernodes: the resulting tree  $T$  has vertex set  $V_T = \{c_k\}$  where  $c_k = \{(j, k)\}_j$  and looking at the resulting edges shows  $T$  is actually an undirected chain  $c_1 - c_2 - \dots - c_h$ . Computationally, since there are  $w = 10$  columns and  $h = 100$  rows, this leads to a reasonably-sized state space ( $2^{10}$  different states) for each supernode.

$$p(x) = \frac{1}{Z} \prod_k \tilde{\psi}_k(x_{c_k}) \prod_k \tilde{\psi}_{k,k+1}(x_{c_k}, x_{c_{k+1}})$$

where

$$\begin{aligned} \tilde{\psi}_k(x_{c_k}) &= \prod_{j=1}^w \psi_{(j,k)}(x_{(j,k)}) \prod_{j=1}^{w-1} \psi_{(j,k),(j+1,k)}(x_{(j,k)}, x_{(j+1,k)}) \\ \tilde{\psi}_{k,k+1}(x_{c_k}, x_{c_{k+1}}) &= \prod_{j=1}^w \psi_{(j,k),(j,k+1)}(x_{(j,k)}, x_{(j,k+1)}) \end{aligned}$$